

# ARIA Evaluator

*Ship safe AI agents with confidence*

Seed Round · \$1.5M · June 2026

ariaeval.io

CONFIDENTIAL

# The Problem

AI agents are being deployed without adequate safety testing

**73%**

of enterprises report AI incidents in the first 6 months of deployment

**\$4.2M**

average cost of an AI-related safety incident (Gartner 2025)

**6–12 mo**

typical model risk validation cycle — manual, qualitative, and too slow for AI agents

**0**

standardised tools exist for adversarial testing of AI agents at scale

# The Solution

ARIA Evaluator: Automated AI agent testing, safety scoring & model risk validation

## 149+ Adversarial Scenarios

14 categories including prompt injection, PII leakage, bias detection, RAG poisoning

## Automated Scoring

AI-powered safety, accuracy, and helpfulness scoring with quality badges

## Model Risk Validation

Quantitative & qualitative risk measures — accelerate MRM cycles from months to days

## Real-Time Transcripts

Watch agent conversations live, replay and analyse failures

## Multi-Model Support

Test across Claude, GPT-4, Llama, Mistral and custom models

## Compliance Evidence

Auto-generated audit trails, scoring reports and risk documentation for regulators

# How It Works

**01**

**Connect**

Point ARIA at your AI agent endpoint — API key or SDK

**02**

**Configure**

Select adversarial scenario categories and test parameters

**03**

**Evaluate**

ARIA runs automated conversations and scores each response

**04**

**Report**

Get detailed safety reports with pass/fail per category

# Market Opportunity



AI Safety & Testing  
(2028 projected)



AI Agent Evaluation  
Tools & Services



Developer-focused  
SaaS evaluation tools

# Product Overview

*Fully built and functional — not a prototype*

## Dashboard

Real-time evaluation metrics, quality badges, plan usage tracking

## Scenario Library

149 pre-built adversarial tests across 14 categories

## Live Transcripts

Watch AI agent conversations in real-time with scoring

## Reports

Detailed safety reports with pass/fail and score breakdowns

## Multi-Region

8 AWS regions — deploy evaluations close to your users

## Compliance

GDPR compliant, SOC 2 ready, 28 compliance documents

# Business Model

PLG motion: free tier → self-serve → enterprise

| Free                 | Individual      | Starter         | Pro             | Unlimited     |
|----------------------|-----------------|-----------------|-----------------|---------------|
| <b>\$0</b>           | <b>\$49/mo</b>  | <b>\$299/mo</b> | <b>\$799/mo</b> | <b>Custom</b> |
| Developers exploring | Solo developers | Growing teams   | Full-scale eval | Enterprise    |

**Net Revenue Retention:** 120%+ target | **Gross Margin:** 80%+ | **CAC Payback:** < 12 months

# Traction & Roadmap

| Done                                                                                                                                                                                      | Q3 2026                                                                                                                                             | Q4 2026                                                                                                                                            | 2027                                                                                                                             |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>• Full product built</li><li>• 149 adversarial scenarios</li><li>• GDPR compliant</li><li>• ariaeval.io live</li><li>• 28 compliance docs</li></ul> | <ul style="list-style-type: none"><li>• Launch free tier</li><li>• First 50 users</li><li>• Design partner LOIs</li><li>• SOC 2 readiness</li></ul> | <ul style="list-style-type: none"><li>• Paid tier launch</li><li>• First \$10K MRR</li><li>• Hire 2 engineers</li><li>• Enterprise pilot</li></ul> | <ul style="list-style-type: none"><li>• \$100K+ MRR</li><li>• SOC 2 Type I</li><li>• Series A ready</li><li>• 15+ team</li></ul> |



# Why ARIA Wins

*\$1.9B+ market (2025) growing 30%+ CAGR — most competitors focus on observability, not offense*

| Capability               | ARIA                       | Maxim AI        | Galileo          | LangSmith      | Arize AI     |
|--------------------------|----------------------------|-----------------|------------------|----------------|--------------|
| Adversarial Red-Teaming  | 149 built-in scenarios     | Limited         | Basic guardrails | None           | None         |
| Proactive Attack Testing | ✓ Core focus               | Simulation only | X                | X              | X            |
| Model Risk Validation    | Quantitative + Qualitative | X               | Partial          | X              | Drift only   |
| Agent Observability      | Cost • Trends • Dimensions | ✓ Strong        | ✓ Strong         | ✓ Core focus   | ✓ Core focus |
| Multi-Model Support      | 20+ models                 | Multi-model     | Multi-model      | LangChain tied | Any model    |
| Compliance Evidence      | GDPR + SOC 2 auto-reports  | Varies          | Varies           | None           | SOC 2        |
| Time to First Test       | 3 minutes                  | ~30 min         | ~1 hour          | ~1 hour        | SDK setup    |

**ARIA's edge:** Offense over defense. Others watch what happened — ARIA tests what could go wrong before you ship.

# The Team

*AI-native founding team with deep platform engineering experience*



## [Founder 1]

CEO & Product

10+ years building developer tools  
Previously: [Company]  
Expertise: AI/ML, Platform Engineering



## [Founder 2]

CTO & Engineering

10+ years in cloud infrastructure  
Previously: [Company]  
Expertise: AWS, Security, SRE

# The Ask

# \$1.5M

Seed round on SAFE note · \$8M post-money cap

Engineering (50%)

\$750K

Sales & Marketing (25%)

\$375K

Infrastructure (15%)

\$225K

Legal & Ops (10%)

\$150K

18 months runway to \$100K+ MRR

# Thank You

*Let's make AI agents safe, together.*

---

**ariaeval.io**  
hello@ariaeval.io